

Practical 7: Introduction to Neural Networks

(Version 1.3)

1 Overview

This practical provides a hands-on introduction to neural networks. The idea is that you will try your hand at *training* a few different types of neural networks and see how well they do. The variations between the networks might depend on several things:

- Choice of network architecture (e.g., the number of hidden nodes)
- Choice of threshold function (also called “transfer function”).
- Choice of learning rate
- Use, or otherwise, of momentum

You will work through some examples using `scikit-learn`, one using the venerable Iris dataset, another on a dataset from the UCI repository, and then you will try writing your own code to learn and then use simple neural networks.

2 A Neural Network for the Iris dataset

We start by looking at a very simple application of the the `scikit-learn` implementation of neural networks.

From KEATS download the file:

```
classify-iris-simple-nn.py
```

and take a look at it in your favourite editor. It should look familiar. That’s because it is pretty much the same as the file that uses a decision tree on the Iris example. Aside from the comments, the only lines which are different are these:

```
iris_nn= MLPClassifier(solver='lbfgs', activation='logistic',  
                      hidden_layer_sizes=(5, 2), random_state=1)  
iris_nn.fit(X, y)
```

and even they look pretty similar to some of the commands we used with decision trees.

The first line creates an instance of the class `MLPClassifier`, for “Multi-Layer Perceptron”, and the second line trains it on the dataset.




Figure 1: The output of a neural network trained on the iris dataset.

As you can see, you have some control over the neural network.¹ What you see here picks a logistic function as the activation function for the neurons, picks the LBFGS solver (the way that the weights are chosen²), fixes the seed of the random number generator (so that results are the same every time), and specifies two hidden layers, one with 5 neurons and one with 2.

Figure 1 shows how well this neural network performs.

Now:

1. Try using neural networks with different architectures (different numbers of neurons in the hidden layers, different numbers of hidden layers). How do these perform?
2. Try using the `sgd` solver (stochastic gradient descent). For that you will likely have to play with some of the learning rate parameters.

3 Classifying Seeds

Now you get to create your own neural network classifier using `scikit-learn`. Download:

`seeds_dataset.csv`

from KEATS. This is (a cleaned up version of) the seeds dataset from the UCI repository:

<https://archive.ics.uci.edu/ml/datasets/seeds>

Also download:

`classify-seeds-starter.py`

¹As ever, there is a manual page for the class, http://scikit-learn.org/stable/modules/generated/sklearn.neural_network.MLPClassifier.html, and this specifies some other parameters that can be set.

²LBFGS “is a solver that approximates the Hessian matrix which represents the second-order partial derivative of a function. Further it approximates the inverse of the Hessian matrix to perform parameter updates. The implementation uses the Scipy version of L-BFGS.” according to http://scikit-learn.org/stable/modules/neural_networks_supervised.html. It is a solver that works well on small data sets.

from KEATS. This loads the data from `seeds_dataset.csv`. Now:

1. Experiment to find a set of features which looks like it will allow a reasonable classification.
Hint: If you plot the data from `seeds_dataset.csv` against pairs of feature values, you can look for groups of seeds that look as easy to separate as the Irises in the Iris dataset.
2. Experiment to find a neural network architecture that can produce a reasonable classification.
Hint: Write code that tests different network architectures and evaluate them using cross-validation. For comparison, the best I could do was a classifier with an accuracy of 0.95.

4 Perceptrons for the Iris Dataset

Now it is time to dig into your own implementation of various approaches to training neural networks. We'll start with a simple perceptron. Starting with another copy of

```
classify-iris-simple-nn.py
```

write code to train a single perceptron to classify the iris data using the error correction procedure. (That is the simplest procedure, so makes a good place to start.) Since you have a single perceptron, you won't be able to do three class classification. But you can classify one class from the other two. Start by learning weights to distinguish class 0 (0 is the value of *y*) from the other two. (This is *setosa*, the very distinct class that appears at the bottom of the plots). When you have cracked that, look to learn weights to distinguish class 2. (This is *virginica*, which appears at the top of the plots.)

Now write code to train a single perceptron using the generalised delta rule. (I started from a fresh copy of `classify-iris-simple-nn.py`.) Again, start by training it to classify class 0 from classes 1 and 2. Then train it to classify class 2 from classes 0 and 1.

Finally, build a single-layer multi-unit neural network that can classify the Iris dataset. If you train one perceptron to classify class 0 from classes 1 and 2, and another to classify class 2 from classes 0 and 1, you can combine the output to classify into three classes. I managed to train mine to an accuracy of 96%.

5 Backpropagation

You can guess what you need to do here. Write a multilayer neural network to classify the Iris data and then train it using backpropagation.

Two things to note here:

1. Since we can classify the Iris dataset pretty well with a single layer network (as in the previous exercise) using a multilayer network is clearly overkill. However, it is a useful exercise to learn how to do backpropagation.
2. Start by using something simple. It might be tempting to build a network with lots of hidden layers and lots of hidden units, but I would suggest something small.

Having said that, my experiments suggest that you probably need at least two output units . . . which also means I'll be impressed if anyone manages to get good results with less than two output units.

6 More

1. For a full set of training procedures, tackle the Iris dataset with a single-layer multi-unit network that is trained using the (non-generalised) Delta rule.
2. See how far you can get in classifying the seeds dataset using a single-layer neural network.
3. Now try a build-your-own neural network on the seeds dataset.

7 Version list

- Version 1.0, February 25th 2019.
- Version 1.1, February 28th 2019.
- Version 1.2, March 1st 2020.
- Version 1.3, February 9th 2021.